

call.

This example would not work for call-by-need. Another example, which would be fine for call-by-value but would hang for call-by-name AND call-by-need:

```
bool __is_foo_called = false; //global state
int bar() {
    if (__is_foo_called) {
        while (true) { ... }
    }
    return 1;
}

void foo(int a) {
    __is_foo_called = true;
    print("%d", a);
}

foo(bar());
```

Here, call-by-value would succeed, because bar is evaluated before any part of the body of foo. In call-by-name and call-by-need, foo sets the global state to true before its parameter is accessed. Then, once its parameter is accessed, bar is evaluated and goes into a loop.

7. Write implementations of the following Prolog predicates. Use a clear and simple style. Do not assume that the input arguments are ground terms. You may assume the standard Prolog predicates like member/2 and append/3, but avoid them in favor of unification if you can. Define auxiliary predicates as needed.

(Note: you can do what you like with the empty list)

7a (5 minutes). `shift_left(L, R)` succeeds if R is the result of "shifting left" the list L by 1. The leading element of L is lost. For example, `shift_left([a,b,c], [b,c])`.

`shift_left([], R)`.

7b (5 minutes). `shift_right(L, R)` is similar, except it shifts right. For example, `shift_right([a,b,c], [a,b])`.

`shift_right([], [])`.
`shift_right([H|L], [H|R]) :- shift_right(L, R)`.

7c (5 minutes). `shift_left_circular(L, R)` is like `shift_left`, except the leading element of L is reintroduced at the right. For example, `shift_left_circular([a,b,c], [b,c,a])`.

`shift_left_circular([H|L], R) :- append(L, [H], R)`.

5 of 10

3/12/19, 9:15 AM

7 of 10

3/12/19, 9:15 AM

```
--or--
shift_left_circular([H|L], R) :- slc(L, H, R).
slc([], X, [X]).
slc([Y|L], H, [Y|R]) :- slc(L, H, R).
```

7d (5 minutes). `shift_right_circular(L, R)` is similar, except it shifts right. For example, `shift_right_circular([a,b,c], [c,a,b])`.

`shift_right_circular(L, R) :- shift_left_circular(R, L)`.

7e (5 minutes). Which of your implementation of the predicates (a)-(d) might be nondeterministic? For each such predicate, give an example call that succeeds more than once.

There will be nondeterminism with any predicate which has multiple rules, because the unification process can choose either and backtrack to try another. So (b), (c), and (d) will all have nondeterminism, because (b) has multiple rules, `append/3` in (c) has multiple rules, and (d) is defined in terms of (c).

8 (10 minutes). What is the set of possible behaviors for each of the following Scheme top-level expressions? Assume that each one is typed in to a fresh instance of a Scheme interpreter. Justify any tricky or obscure parts of your answers by citing particular parts of the formal semantics of Scheme.

(not (if #f #f))
 You are not responsible for knowing the two-argument if statement, but FYI it will nondeterministically return #t or #f if the conditional is #f. So this can either evaluate to #f or #t.

(car (begin (set! car cdr) (list car cdr)))
 See Dybvig section 2.3 for comments on order of evaluation of Scheme expressions. To summarize, different implementations may evaluate in different orders, possibly right-to-left or left-to-right. The value of 'car' may be looked up to be the 'car' function before the right expression (begin ...) is evaluated, or the begin expression may be evaluated first.

In the first case (left-to-right), the first 'car' would be evaluated to the 'car' function, and then the 'car' identifier would be set to the 'cdr' function. Then, the list of the 'car' identifier (looked up to be 'cdr') and the 'cdr' function would be created. The 'car' from the beginning would access the first element, which is the 'cdr' function (from the 'car' identifier).

In the second case (right-to-left), the (list...) subexpression would get evaluated first, created a list of the 'car' function and the 'cdr' function. Then the 'car' identifier would get bound to the 'cdr' function, and finally the leftmost 'car' identifier would be looked up to be the 'cdr' function. So (cdr (car cdr)) would evaluate to the 'cdr' function.

6 of 10

3/12/19, 9:15 AM

Note, however, that it is not even guaranteed that if an expression is evaluated left-to-right, its subexpressions will also be evaluated left-to-right. Directionality may change in subexpressions, so you may get some combination of the above behavior.

9 (5 minutes). The abstract syntax for Scheme in R5RS section 7.2.1 lists only constants, identifiers, procedure calls, lambda, if, and set!. The following expressions aren't any of these things, so why are they valid expressions in Scheme? Or if they're not valid expressions, say why not.

```
(let ((a '())) a)
(cond ((eq? '() '()) 1) (#t 0))
```

The above two are macros that are replaced with the abstract syntax defined in R5RS.

(1) - This is not valid because 1 is not a function.

10 (20 minutes). Consider the following Java code snippet, taken from BioJava. Translate it as best you can into idiomatic Python. For parts that you cannot easily translate, say what the problems are, and how you'd go about addressing these problems if you actually had to translate all of BioJava into Python.

```
package org.biojava.bio.program.ensembl;
import java.util.*;
import java.lang.ref.*;
import org.biojava.bio.*;

...
class CloneDB extends ImmutableSequenceDBBase
    implements SequenceDB {
    private Ensembl ensembl;

    private Map cloneCache;
    private Set ids_cache;

    { cloneCache = new HashMap(); }

    public CloneDB(Ensembl db) {
        this.ensembl = db;
    }

    public String getName() {
        return "http://www.ensembl.org/";
    }

    Sequence getCloneSequence(Clone cont)
        throws EnsemblException
    {
        Sequence seq = (Sequence)
            cloneCache.get(cont.getID());
        if (seq == null) {
            if (ensembl.getUseHeavyClones()) {
                seq = new EnsemblHeavyCloneSequence(
                    cont, ensembl);
            } else {
                seq = new EnsemblCloneSequence(
                    cont, ensembl);
            }
            cloneCache.put(cont.getID(), seq);
        }
        return seq;
    }
    ...
}
```

(Note: you will not be responsible for knowing this exact problem, since it covers a HW which is no longer assigned. However, you should be able to convert Java code to Python and comment on the limitations.)

```
from java.util import *
from java.lang.ref import *
from org.biojava.bio import *
```

```
class CloneDB(ImmutableSequenceDBBase):
    def __init__(self, db):
        self.ensembl = db
        self.cloneCache = HashMap()

    def getName(self):
        return "http://www.ensembl.org/"

    def getCloneSequence(self, cont):
        seq = self.cloneCache.get(cont.getID())
        if seq == None:
            if self.ensembl.getUseHeavyClones():
                seq = EnsemblHeavyCloneSequence(cont, self.ensembl)
            else:
                seq = EnsemblCloneSequence(cont, self.ensembl)
            self.cloneCache.put(cont.getID(), seq)
        return seq
```

There are many issues.

1. There is not an equivalent to the package keyword. To achieve the same package semantics, you have to use directory-local `__init__.py` files as described here: <https://docs.python.org/2/tutorial/modules.html#packages>
2. There are no interfaces in Python, so you have to assume or manually check that your CloneDB class is implementing the correct functions.
3. Python does not really have a concept of private/public fields and methods. There is a trick you can do, detailed here: <http://stackoverflow.com/questions/70528>

8 of 10

3/12/19, 9:15 AM

4. Type information cannot be translated to Python because Python is dynamically typed.
5. The `{ cloneCache = new HashMap(); }` line is an "initializer block", which is not really an issue, because in Java, the line is simply copied into every constructor by the compiler. It simply allows you to share large code blocks between different constructors. To translate to Python, just stick it in the `__init__` method.
6. Python has no exception annotations, but you throw and catch similar to Java, so the "throws `EnsembleException`" is just left off.
7. You would have to assume that the Java libraries imported have also been ported to Python.

11 (15 minutes). For each of the following seemingly-arbitrary language rules, explain what performance gains (if any) arise because of the rules. Analyze both time and space performance. List the rules roughly in decreasing order of importance, performance-wise.

- * Arrays are origin-zero, not origin 1.

Slightly faster. For origin-zero, to get the address of `a[4]`, you will calculate `a + 4*N`, where `N` is the size of the array cell. For origin-1, to get the address of `a[5]`, you calculate `a + (5-1)*N`, which is extra arithmetic.

- * Arrays must be allocated statically or on the stack; they cannot be allocated on the heap.

This allows you to do away with any runtime memory manager (e.g. garbage collector), because all your memory will be cleaned up when a stack frame is popped off.

- * Each object in an array must have the same size.

Faster, because you can do arithmetic indexing. However, this may waste space, because you have to make the size the maximum of whatever objects you want to put in.

- * Subscript violations result in undefined behavior; an exception may or may not be raised.

Performance gain because you don't have any bounds-checking overhead. You also get a small space gain because you don't have to store the array length anywhere.

- * The array must have just one dimension.

Performance loss due to the extra index calculation of both dimensions. Also the possibility of external fragmentation preventing the allocation

of a large array (which would otherwise be smaller chunks given a C++ or Java way of doing multidimensional arrays).

- * The number of elements in an array is fixed at compile-time.

If you're using stack allocation of arrays, fixed size makes this faster. However, you may have to allocate a large array according to the max size it would become if it was dynamically allocated.

- * Arrays must always be accessed directly; there are no pointers to arrays, and arrays cannot be passed by reference to procedures.

Very space-wasting, because you must copy arrays when passing them into functions. However, this avoids aliasing, which allows for better compiler optimizations, including constant propagation and code reordering.

12 (15 minutes). Suppose we add a new statement "return[N]EXPR;" to our C implementation, as an extension to C. `N` must be an integer constant expression with a nonnegative value, and `EXPR` can be any valid C expression. This new statement acts like "return EXPR;", except that it causes the `N`th caller of the current function to return with the value of the expression `EXPR`. For example, suppose `F` calls `G` which calls `H`, and suppose `H` then executes "return[2]5;". This causes `F` to return immediately, with the value 5; `G` and `H` are silently exited in the process. By definition, "return[0]EXPR;" is equivalent to "return EXPR;".

What problems do you see in implementing "return[N]EXPR;"? Are these problems inherent to C, or can they be worked around by changing the semantics of the new statement slightly? Would it be easier to implement this new statement in some of the other languages that we have studied? If so, which ones and why? If not, why not?

A huge problem with this is centered around static typing. Because some function `F` does not know how deep in the call chain its return value will come from, it cannot know the type of that return value. Imagine the functions `F`, `G`, and `H` in the problem have return types `void`, `int`, and `double`, respectively. `H` executing "return[2]1.0" would cause `F` to try to return 1.0. These problems are inherent to C.

This would be much easier to implement in Python, for example, because of dynamic typing. The return type of the function is not specified, so returning some alternative value would not be a problem.

This sample midterm is the midterm from Spring 2008. If you see any errors in the sample answer, please email me or tell me on Piazza. Good luck on your midterms!

1. "Ireland has leprechauns galore." is an example of a particular kind of syntactic construct in English. Can you construct a similar example in C++, OCaml, or Java? If so, give an example; if not, explain why not

- Artificial languages are generally more carefully-designed than natural languages. Thus, these kinds of exceptions should be rare, if they exist at all.

2.

a) Write an OCaml function 'twice' that accepts a function f and returns a function g such that $g(x)$ equals $f(f(x))$. For simplicity's sake, you can assume that f is free of side effects, and you can impose other restrictions on f and x . Try to keep the restrictions as minor as possible, and explain any restrictions you impose. Or, if 'twice' cannot be written easily in OCaml, explain why not.

- let twice f $x = f(f(x))$
- restriction is that f is 'a' \rightarrow 'a'

b) Same as a) except write a function 'half' that accepts a function f and returns a function g such that $f(x)$ equals $g(g(x))$.

- It is difficult, since its implementation must depend on the specific function f used.
- For example, if f is identity function $f(x) = x$, then g is easy-- also the identity function.
- However, if f is the sine function, then g is super difficult-- given x , $g(g(x)) = \sin(x)$???
- The restriction, like a), is that f is 'a' \rightarrow 'a'

c) Give the types of 'twice' and 'half'

- twice: ('a' \rightarrow 'a') \rightarrow 'a' \rightarrow 'a'
- half: ('a' \rightarrow 'a') \rightarrow 'a' \rightarrow 'a'

3. Consider the following grammar for a subset of the C++ language.

expression:
expression ? expression : expression
expression != expression
expression + expression

! expression
INTEGER-CONSTANT
(expression)

For example, $((! 0 + 1 = 2 ? 3 : 4))$ is read as "if not-not-0 plus 1 does not equal 2, then 3 else 4, and evaluates to 4.

a) What are the tokens of this subset of C++?
? : != + ! INTEGER-CONSTANT ()

b) Show that this grammar is ambiguous
Draw the two parse trees for expression $1+2+3$

c) Rewrite the grammar so that it is no longer ambiguous, resolving any ambiguities in the same way that C++ does. Recall that in C++, the expression

$(0 != 1 != 2 || 3 + 14 + 5 || 6 ? 7 : 8 ? 9 : 10)$ is like

$(((((0 != 1) != 2) || ((3 + (14)) + 5)) || 6) ? 7 : (8 ? 9 : 10))$

$E \rightarrow E2 ? E2 : E \mid E2$
 $E2 \rightarrow E2 != E3 \mid E3$
 $E3 \rightarrow E3 + E4 \mid E4$
 $E4 \rightarrow !E4 \mid E5$
 $E5 \rightarrow \text{INTEGER-CONST} \mid (E)$

d) Translate the rewritten grammar into a syntax diagram

eg. $E2$
0-----> $E3$ -----> 0
[<-- != <--]

4. A numerical analyst is really bothered by the special value of IEEE floating point, and asks you to modify Java C++ to fix what she views as a serious conceptual flaw. She wants her Java programs to throw an exception instead of returning infinities and NaNs. Is her request reasonable for Java C++ programs? Is it implementable? Why/why not? Don't worry about compatibility with existing compilers, etc.; assume that you are the inventor of Java C++ and she is asking for this feature early in your language design process.

Scanned by CamScanner

Scanned by CamScanner

- It's kind of reasonable...maybe the numerical analyst has applications that don't allow Infs and NaNs. However, this gets rid of the choice of having these special values. To implement it, Java C++ can throw an exception whenever it encounters a special value.

5. Give an example of four distinct Java C++ types A, B, C, D such that A is a subtype of B, B is a subtype of C, C is a subtype of D, and C is a subtype of D. Or, if such an example is impossible, explain why not.

- class D {...}
- class A: public B, public C {...}
- class B: public D {...}
- class C: public D {...}

6. Explain how you would implement OCaml-style type checking, in an implementation that uses dynamic linking heavily. What problems do you foresee in programs that relink themselves on the fly?

- (verbatim from the professor) The problem is that one needs to do type checking as well as the usual name checking that link editing normally does. And the type checking needs to follow OCaml's rules. In particular it needs to work with generic types, where the type of the linked-to function does not exactly match the type needed by the caller, but the types will match after applying a unifier.

7.

a) Write a curried OCaml function "interleave C S l1 l2" that constructs a new list L from the lists l1 and l2, using the chooser C with seed S, and returns a pair (S1, L) where S1 is the resulting seed and L is the interleaved list. For example, "interleave C S [1;2] [3;4;5]" might invoke C four times and then return (S1, [1;3;4;2;5]). At each step of the iteration, "interleave" should use the chooser to decide whether to choose the first item of l1 or the first item of l2, when deciding which of the two items to put next into L. If l1 is empty, the chooser need not be used, since "interleave" will just return l2; and similarly if l2 is empty, "interleave" should just return l1 without invoking C.

Chooser C is defined
<http://www.cs.ucla.edu/classes/spring08/cs131/hw/hw1.html>

"interleave" should invoke C a minimal number of times, left to right across the lists l1 and l2. "interleave" should avoid side effects; it should be written in a functional style, without using OCaml libraries.

```
let rec interleave c s l1 l2 =
  if l1 = [] then (s, l2) else if l2 = [] then (s, l1) else
  match (c s) with
  | (true, s1) -> (match l1 with
    | h::t -> (let (s2, l1) = (interleave c s1 t l2) in
      (s2, h::l1)))
  | (false, s1) -> (match l2 with
    | h::t -> (let (s2, l1) = (interleave c s1 l1 t) in
      (s2, h::l1)))
```

b) Write a function "outerleave" that does the opposite of what "interleave" does: it splits a list into two sublists that can be interleaved to get the original list, and returns a triplet consisting of the new seed and the two sublists. That is, "outerleave C S [1;3;4;2;5]" might yield (S1, [1;3;2], [4;5]). If given a list of length N, "outerleave" always invokes the chooser N times.

```
let rec outerleave c s l =
  match l with
  | [] -> (s, [], [])
  | h::t -> match (c s) with
    | (true, s1) -> (let (s2, l1, l2) = (outerleave c s1 t) in
      (s2, h::l1, l2))
    | (false, s1) -> (let (s2, l1, l2) = (outerleave c s1 t) in
      (s2, l1, h::l2))
```

c) Give the data types of all top-level values or functions defined in your answer to a) and b).

- interleave: ('a -> bool * 'a) -> 'a -> 'b list -> 'b list -> 'a * 'b list
- outerleave: ('a -> bool * 'a) -> 'a -> 'b list -> 'a * 'b list * 'b list
- chooser c: 'a -> (bool, 'a)
- seed s: 'a
- list l1/l2: 'b list

Scanned by CamScanner

Scanned by CamScanner

UCLA CS 131 Midterm, Fall 2017
100 minutes total, open book, open notes
closed computer

Name:	Student ID:						
1	2	3	4	5	6	total	

1a (10 minutes). Write an OCaml function `merge_sorted` that merges two sorted lists. Its first argument should be a comparison function 'lt' that compares two list elements and returns true if the first element is less than the second. Its second and third arguments should be the lists to be merged. For example, `(merge_sorted (<) [21; 49; 61] [-5; 20; 25; 49; 50; 100])` should yield `[-5; 20; 21; 25; 49; 49; 50; 61; 100]`.

1b (3 minutes). What is the type of `merge_sorted`?

1c (3 minutes). What does the following expression yield, and what is its type?

```
merge_sorted (fun a b -> List.length a < List.length b)
```

1d (8 minutes). Is your implementation of `merge_sorted` tail-recursive? If so, briefly say why it won't have any problem with stack overflow. If not, briefly say why not, and explain any problems you would have in rewriting your implementation to make it tail-recursive.

2 (9 minutes). Consider the following top-level OCaml definitions:

```
let f f = f 1 1
let g g = g 0.0 g
let h h = h f "x"
```

For each identifier declared in this code, give the identifier's scope and type. Or, if there is a scope or type error, briefly explain the error.

3a (5 minutes). In Java, is the subtype relation transitive? That is, if A is a subtype of B and B is a subtype of C, is A a subtype of C? If so, explain why; if not, give a counterexample.

3b (5 minutes). In Java, is the graph of the subtype relation a tree? If so, explain why, and say what the root is; if not, give a counterexample.

4. Consider the following grammar for declarations in a subset of C.

incorrect for C. Prove that it is syntactically correct. Briefly explain why it is semantically incorrect.

4c (5 minutes). Suppose we changed the grammar by replacing the ruleset for type-qualifier-list with the following:

```
type-qualifier-list:
  type-qualifier type-qualifier-list?
```

Would this cause any problems? If so, describe a problem and give an example. If not, briefly explain why not.

4d (10 minutes). Suppose we changed the original grammar by replacing the two rulesets for declarator and direct-declarator with the following single ruleset:

```
declarator:
  pointer? declarator
  ID
  '(' declarator ')'
  declarator '[' INT '['
  declarator '(' 'void' ')'
```

Would this cause any problems? If so, describe a problem and give an example. If not, briefly explain why not.

4e (10 minutes). Draw a syntax chart for the original grammar.

5 (10 minutes). Suppose we write Java code in a purely functional style, in that we never assign to any variables except when initializing them. That is, we always initialize local variables and never assign to them later, and we always initialize instance variables once at the start of constructors and never assign to them later.

In our purely-functional Java programs, is the Java Memory Model still relevant, or can we ignore it? If it's still relevant, explain which parts of it still apply and give an example. If not, briefly explain why not.

6 (12 minutes). Consider the following code, taken from the answer to the older version of Homework 2.

```
let match_empty frag accept = accept frag

let match_nothing frag accept = None

let rec match_star matcher frag accept =
  match accept frag with
  | None ->
    matcher frag
    (fun frag1 ->
      if frag == frag1
```

This grammar uses a form of EBNF in which the left hand side is not indented and is followed by '+', each right hand alternative is indented, nonterminals are strings of letters and '+', terminal symbols are either surrounded by single quotes or are INT (meaning an integer constant) or ID (meaning an identifier), and X? stands for zero or one instances of X.

```
declaration:
  declaration-specifiers init-declarator-list? '+'
```

```
declaration-specifiers:
  storage-class-specifier declaration-specifiers?
  type-specifier declaration-specifiers?
  type-qualifier declaration-specifiers?
  function-specifier declaration-specifiers?
```

```
storage-class-specifier:
  'typedef'
  'static'
  'char'
  'int'
```

```
type-qualifier:
  'const'
  'volatile'
  'inline'
  '_Noreturn'
```

```
init-declarator-list:
  init-declarator
  init-declarator-list ',' init-declarator
```

```
init-declarator:
  declarator
  declarator '=' initializer
```

```
declarator:
  pointer? direct-declarator
```

```
direct-declarator:
  ID
  '(' declarator ')'
  direct-declarator '[' INT '['
  direct-declarator '(' 'void' ')'
```

```
pointer:
  '*' type-qualifier-list? pointer?
```

```
type-qualifier-list:
  type-qualifier-list? type-qualifier
```

```
initializer:
  ID
  INT
```

4a (2 minutes). What makes this grammar EBNF and not simply BNF?

4b (8 minutes). Give an example declaration that is syntactically correct (i.e., it is produced by this grammar) but is semantically

```
then None
else match_star matcher frag1 accept)

| ok -> ok

let match_nucleotide nt frag accept =
  match frag with
  | [] -> None
  | n::tail -> if n == nt then accept tail else None

let append_matchers matcher1 matcher2 frag accept =
  matcher1 frag (fun frag1 -> matcher2 frag1 accept)

let make_appended_matchers make_a_matcher ls =
  let rec mams = function
  | [] -> match_empty
  | head::tail -> append_matchers (make_a_matcher head) (mams tail)
  in mams ls
```

In this code, a matcher is a curried function taking two arguments: first, a fragment 'frag' and second, an acceptor 'accept'. Suppose we change the API for matchers by interchanging their arguments, so that the acceptor comes first (all the functions remain curried). Rewrite the above code to use the altered API, and simplify the resulting code as much as possible.